

Assignment3

MengZhang

2019/3/15

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
redwine = read.csv("redwine.csv")
redwine$final_quality <- with(ifelse(quality>mean(quality), 1, 0), data=redwine)
set.seed(1)
glm.fit <- glm(final_quality ~ .-quality, data=redwine,family="binomial")
summary(glm.fit)
```

```
##
## Call:
## glm(formula = final_quality ~ . - quality, family = "binomial",
##      data = redwine)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4025  -0.8387   0.3105   0.8300   2.3142
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    42.949948   79.473979   0.540  0.58890
## fixed.acidity     0.135980   0.098483   1.381  0.16736
## volatile.acidity  -3.281694   0.488214  -6.722 1.79e-11 ***
## citric.acid      -1.274347   0.562730  -2.265  0.02354 *
## residual.sugar     0.055326   0.053770   1.029  0.30351
## chlorides        -3.915713   1.569298  -2.495  0.01259 *
## free.sulfur.dioxide  0.022220   0.008236   2.698  0.00698 **
## total.sulfur.dioxide -0.016394   0.002882  -5.688 1.29e-08 ***
## density          -50.932385  81.148745  -0.628  0.53024
## pH               -0.380608   0.720203  -0.528  0.59717
## sulphates         2.795107   0.452184   6.181 6.36e-10 ***
## alcohol          0.866822   0.104190   8.320 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2209.0  on 1598  degrees of freedom
## Residual deviance: 1655.6  on 1587  degrees of freedom
## AIC: 1679.6
##
## Number of Fisher Scoring iterations: 4
```

```

train <- sample(dim(redwine)[1], dim(redwine)[1] / 2)

fit.glm <- glm(final_quality ~ .-quality, data = redwine, family = "binomial", subset = train)
summary(fit.glm)

##
## Call:
## glm(formula = final_quality ~ . - quality, family = "binomial",
##      data = redwine, subset = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0723  -0.9083   0.3579   0.8787   2.1113
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -59.614569  103.475941  -0.576  0.564534
## fixed.acidity     -0.014417   0.127280  -0.113  0.909816
## volatile.acidity  -2.185291   0.642672  -3.400  0.000673 ***
## citric.acid       -0.967160   0.788819  -1.226  0.220166
## residual.sugar    -0.010555   0.070520  -0.150  0.881017
## chlorides         -6.702735   2.234080  -3.000  0.002698 **
## free.sulfur.dioxide  0.022355   0.011445   1.953  0.050786 .
## total.sulfur.dioxide -0.017249   0.004164  -4.142  3.44e-05 ***
## density          58.890929  105.587445   0.558  0.577018
## pH               -2.139387   0.973935  -2.197  0.028046 *
## sulphates         3.236184   0.663447   4.878  1.07e-06 ***
## alcohol          0.827697   0.137117   6.036  1.58e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1106.12  on 798  degrees of freedom
## Residual deviance:  868.78  on 787  degrees of freedom
## AIC: 892.78
##
## Number of Fisher Scoring iterations: 4

probs <- predict(fit.glm, newdata = redwine[-train, ], type = "response")
pred.glm <- rep(0, length(probs))
pred.glm[probs > 0.5] <- 1

mean(pred.glm != redwine[-train, ]$final_quality)

## [1] 0.23125

boot.fn <- function(data, index) {
  fit <- glm(final_quality ~ .-quality, data = data, family = "binomial", subset = index)
  return (coef(fit))
}
library(boot)
boot(redwine, boot.fn, 1000)

##
## ORDINARY NONPARAMETRIC BOOTSTRAP

```

```
##
##
## Call:
## boot(data = redwine, statistic = boot.fn, R = 1000)
##
##
## Bootstrap Statistics :
##      original      bias      std. error
## t1*   42.94994813 -4.4614677337 84.419963130
## t2*    0.13598034 -0.0025725162 0.103200422
## t3*   -3.28169367 -0.0582294814 0.529453903
## t4*   -1.27434734 -0.0461774161 0.589334738
## t5*    0.05532602 -0.0037607448 0.065942942
## t6*   -3.91571291 -0.1295325690 1.788169423
## t7*    0.02222037 0.0002795740 0.008217903
## t8*   -0.01639392 -0.0002373653 0.003022671
## t9*  -50.93238519  4.5368754895 86.128740298
## t10* -0.38060751 -0.0182062730 0.737308577
## t11*  2.79510651 0.0724327547 0.547146073
## t12*  0.86682223 0.0042059078 0.107827776

fit.glm <- glm(final_quality ~ .-quality, data = redwine, family = "binomial")
summary(fit.glm)

##
## Call:
## glm(formula = final_quality ~ . - quality, family = "binomial",
##      data = redwine)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4025  -0.8387   0.3105   0.8300   2.3142
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   42.949948  79.473979   0.540  0.58890
## fixed.acidity    0.135980   0.098483   1.381  0.16736
## volatile.acidity -3.281694   0.488214  -6.722 1.79e-11 ***
## citric.acid     -1.274347   0.562730  -2.265  0.02354 *
## residual.sugar   0.055326   0.053770   1.029  0.30351
## chlorides       -3.915713   1.569298  -2.495  0.01259 *
## free.sulfur.dioxide 0.022220   0.008236   2.698  0.00698 **
## total.sulfur.dioxide -0.016394   0.002882  -5.688 1.29e-08 ***
## density        -50.932385  81.148745  -0.628  0.53024
## pH              -0.380608   0.720203  -0.528  0.59717
## sulphates       2.795107   0.452184   6.181 6.36e-10 ***
## alcohol         0.866822   0.104190   8.320 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 2209.0  on 1598  degrees of freedom
## Residual deviance: 1655.6  on 1587  degrees of freedom
## AIC: 1679.6
```

```
##  
## Number of Fisher Scoring iterations: 4
```